Contribution statement

Authors: Emmanuel Azubuike (2698752), George Ayensu (2722188), Banele Mkhize (2663814).

All three authors contributed equally to this project, one third each. Every author took part in each stage below, and all design and interpretive decisions were made jointly.

1. **Problem formulation and background.** Reading the archive paper and related work on these newspapers, agenda-setting and framing, embedding-based change detection and tokenizer fairness; formulating the research question and hypotheses H1–H3; identifying the main threats to a fair comparison (different years, corpus size, word frequency, change over time, one paper per side) and designing a response to each.

2. **Data preparation and exploration.** Loading and cleaning the transcriptions of both newspapers; building date-matched, size-matched samples; exploratory analysis (a grouped cross-validated classifier and yearly rate series) that led to the religion concept; building and checking the concept word stems, including correcting the *law* stem.

3. **Salience analysis (H1).** Rate ratios with issue-bootstrap confidence intervals; frequency-matched random word sets as the null; within-paper era drift; detecting readers' letters and checking the detector by hand; comparing letters with the rest of each paper; the same-editor and same-control contrasts that separate control from editor and time.

4. **Framing analysis (H2).** The hyper-parameter search for stable embeddings on held-out years; profile-shift and neighbour-overlap measures; the permutation null; the directional anchor test; robustness checks (wider windows, joined word forms, removing *land* and the anchor *endlu*); the detection floor; the word-sense analysis of *umhlaba*.

5. **Tokenizer fairness (H3).** Measuring subword cost for XLM-R, mBERT and Serengeti on modern, spelling-altered and historical text, and analysing the tokenizers' vocabularies.

6. **Responsible NLP and ethics.** Assessing the limits and risks of interpretation (speaker review, extractive readings, spelling normalisation, archive labelling); completing the NeurIPS checklist and the Faculty AI declaration.

7. **Writing and reproducibility.** Designing the figures and tables; writing the extended abstract and supplementary material and revising them through several rounds of review; organising the code, `README.txt` and `requirements.txt`; checking that every reported number is regenerated by the submitted code.

All authors reviewed and approved the final report and code, and take joint responsibility for them.